

Computational Statistics

Lab 3

Emil K Svensson and Rasmus Holm

2017-02-16

Question 1

1.1

```
pop <- read.csv2("../data/population.csv", encoding = "latin1")
```

1.2

```
cityUnif <- function(data){  
  data$prob <- data$Population / sum(data$Population)  
  data$cumprob <- cumsum(data$prob)  
  return(which.min(data$cumprob < runif(1)))  
}
```

1.3

```
rmpop <- pop  
  
set.seed(123456)  
while(nrow(rmpop) > 20){  
  rmpop <- rmpop[-cityUnif(rmpop),]  
}
```

1.4

```
print(paste(as.character(rmpop$Municipality), ":", rmpop$Population))
```

```
## [1] "Vingåker : 8911"      "Ydre : 3672"  
## [3] "Aneby : 6446"        "Borgholm : 10806"  
## [5] "Dals-Ed : 4729"      "Grästorp : 5857"  
## [7] "Gullspång : 5335"    "Hjo : 8859"  
## [9] "Strömstad : 11690"   "Hällefors : 7333"  
## [11] "Ljusnarsberg : 5055" "Skinnskatteberg : 4567"  
## [13] "Rättvik : 10797"     "Bräcke : 6865"  
## [15] "Dorotea : 2900"      "Vilhelmina : 7156"  
## [17] "Arjeplog : 3143"     "Jokkmokk : 5210"  
## [19] "Älvsbyn : 8387"      "Övertorneå : 4920"
```

As seen in the output above, we in general select cities with small populations. This is logical since the removal of cities is based on the population, i.e. highly populated cities are more likely to be removed, thus it is more likely to leave smaller cities in the output.

1.5


```
library(ggplot2)

true_pop <- data.frame(pop=pop$Population)
picked_pop <- data.frame(pop=rmpop$Population)

true_pop$data <- "True"
picked_pop$data <- "Selected"

plot_data <- rbind(true_pop, picked_pop)

ggplot(plot_data, aes(pop, fill=data)) +
  geom_histogram(color="black", binwidth=10000) +
  facet_grid(data ~ ., scales="free_y") +
  xlab("Population") + ylab("Count") +
  labs(fill="Data")
```




In the histograms above the width of a staple is 10000 and we can see that in the data set most cities are ≤ 50000 in population and the cities we selected are all ≤ 20000 in size. This is as described above expected but we can also see that it is expected in terms of the number of cities in that range.

Question 2

2.1

```
invLap <- function(){  
  rng <- runif(1)  
  
  if(rng > 0.5) {  
    return(-log(2 - 2 * rng) )  
  } else {  
    return(log(2 * rng))  
  }  
}  
  
plotdata <- sapply(1:10000,FUN = function(x) invLap())  
  
hist(plotdata, prob = TRUE, ylim = c(0,0.5), xlim = c(-6,6),  
     xlab = "x", main = "Laplace distribution ~(0,1)",  
     sub = "Inverse CDF method", breaks = 20)  
lines(density(plotdata), col = "red", lwd = 2)
```

Laplace distribution $\sim(0,1)$



Yes, the histogram and the density-curve we included seems reasonable compared to the normal shape of a Laplace distribution.

2.2

```
DE01 <- function(x) {
  (1 / 2) * exp(-abs(x))
}

ar <- function(c) {
  rej <- 0

  generated <- FALSE
  x <- c()

  while(!generated){
    y <- invLap() #  $Y \sim f_y$ 
    u <- runif(1) #  $U(0,1)$ 

    fx <- dnorm(y, mean = 0, sd = 1) #  $f_x(y)$ 
    fy <- DE01(y) #  $f_y$ 

    if(u < fx / (c * fy)){
      return(c(y, rej)) #alt. set x <- y and generated = TRUE, to end the loop
    }


    rej <- rej + 1
  }
}
```

Here is our function.

Choosing the appropriate c

```
set.seed(123456)
n <- 100000
lap <- sapply(1:n, FUN = function(x) invLap())
nor <- rnorm(n, 0, 1)

plot(x = 0, y = 0, col = "white", xlim = c(-6,6),
     ylim = c(-0.01, 0.5), xlab = "x",
     ylab = "Density" )
lines(density(lap), col = "red", lwd = 2)
lines(density(nor), col = "blue", lwd = 2)
```



Our approach for approximating c was to try to make a laplace-curve that would as much as possible cover the whole density-curve for the normal distribution.

For $x > 0$ since it is symmetric

$$cf_y(x) - f_x(x) = 0$$

$$c/2e^{-x} - 1/(\sqrt{2\pi})e^{-x^2/2} = 0$$

$$\log(c/2) - x - \log(1/(\sqrt{2\pi})) + x^2/2 = 0$$

after some equation-solving we come to the solution


$$c = 1.315$$

Simulation and comparison

```
ARsim <- data.frame(rn = 1, rej = 1)
c <- 1.32

for(i in 1:2000){
  ARsim[i,] <- ar(c=c)
}


plot(x = 0, y = 0, col = "white", xlim = c(-5,5),
      ylim = c(-0.01,0.5), xlab = "x", ylab = "Density" )
lines(density(ARsim$rn), col="red", lwd=2)
lines(density(nor), col="blue", lwd=2)
```




```
par(mfrow = c(1,2))
hist(ARsim$rn, main = "Simulated Normal", xlab = "x",
     col = "orange", breaks = 30)

set.seed(123456)
hist(rnorm(2000,0,1), main = "Normal drawn from rnorm", xlab = "x",
     col = "lightgreen", breaks = 30 )
```

Simulated Normal



Normal drawn from rnorm



The simulated one is a bit more rough around the edges and more narrow in the center in comparison to the

histogram drawn from the rnorm.

The expected rejection rate is $\frac{n(M-1)}{nM}$ since the expected number of iterations before a accepted value is M so intuitively the number of rejections is $M - 1$. n here is only the number of values that we use, but these cancel each other out so they are not necessary but we left them here for clarity. The theoretical rejection rate would be 0.242 and what we empirically found was 0.262 which is reasonably close.